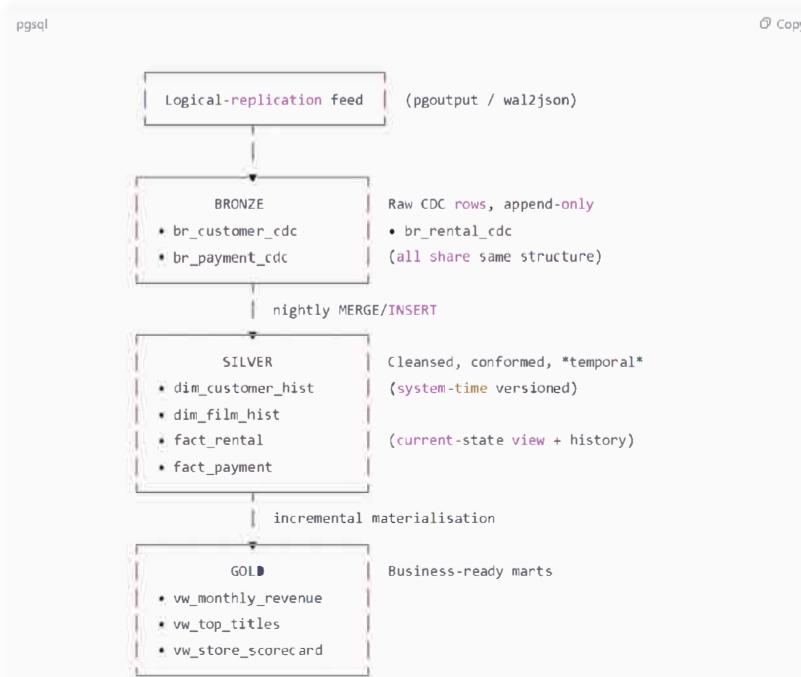


Below is a hands-on **medallion blueprint** that turns the *Pagila* OLTP schema into a three-layer lakehouse while showcasing PostgreSQL-style **temporal tables**.

Everything fits in 9 **physical tables + 3 convenience views**, so the repo stays demo-friendly.

1 High-level flow



2 Bronze layer – raw change capture

Table	Technique	Why it's useful
br_customer_cdc	CREATE TABLE ... (tsn bigint, op char(1), full_row jsonb, load_ts timestamptz)	Preserves every insert/update/delete with the exact payload.
br_rental_cdc	Same columns	High-volume fact → perfect for watermark demos.
br_payment_cdc	Same columns (even though <i>Pagila</i> already partitions by month)	Shows that <i>CDC ≠ partition strategy</i> and keeps the bronze schema uniform.

Implementation hint → subscribe to WAL using `wal2json` and land each message verbatim. You now have a temporal “transaction-time” log without writing a line of JSON parsing logic.

3 Silver layer – conformed & system-versioned

3.1 Temporal dimensions

Create *system-time* tables that keep `valid-from` / `valid-to` automatically:

```
sql Copy  
CREATE TABLE dim_customer_hist (
    customer_id      int,
    first_name       text,
    last_name        text,
    email            text,
    active           bool,
    sys_start        timestampz NOT NULL DEFAULT clock_timestamp(),
    sys_end          timestampz NOT NULL DEFAULT 'infinity',
    PRIMARY KEY(customer_id, sys_start)
) WITH (system_versioning = true); -- PG 15+ extension / trigger template
```

Replace straight `INSERT / UPDATE` with:

```
sql Copy  
MERGE INTO dim_customer_hist tgt
USING staging_customer src
ON   tgt.customer_id = src.customer_id
WHEN MATCHED AND (src.hash <> tgt.hash)    -- change detected
    THEN UPDATE SET ... , sys_end = clock_timestamp()
WHEN NOT MATCHED
    THEN INSERT (cols...) VALUES (...);
```

You now get *built-in* “as-of” querying:

```
sql Copy  
SELECT * FROM dim_customer_hist
FOR SYSTEM_TIME AS OF '2023-07-01'          -- what the warehouse "looked like"
WHERE customer_id = 42;
```

3.2 Current-state views

```
sql Copy  
CREATE OR REPLACE VIEW dim_customer AS
SELECT * EXCEPT (sys_end)
FROM   dim_customer_hist
WHERE  sys_end = 'infinity';
```

3.3 Facts

Rental and *payment* facts rarely change after they close.

Load them once, add a *partition key* (`rental_date` / `payment_date`) and optionally a *point-in-time snapshot* column to prove the temporal model.

4 Gold layer – business marts & KPIs

Mart/view	Grain	Sample metric
vw_monthly_revenue	store-month	<code>SUM(amount)</code>
vw_top_titles	film-month	rentals, revenue, average days out
vw_store_scorecard	store-day	revenue, late-return %, active customers

All three read **only current rows** from `dim_*` views plus the two fact tables, so refresh is cheap and safe.

5 Why temporal tables *belong* in medallion

Stage	Temporal purpose
Bronze	Already is a transaction-time history (raw WAL payload).
Silver	Turn <i>transaction time</i> into <i>application time</i> (validity). System-versioned tables automate SCD Type-2, keep the model elegant and make time-travel demos fun.
Gold	Mostly current-state, but analysts can trivially change <code>dim_customer</code> → <code>dim_customer_hist</code> to answer “what-if we reran the report last quarter?”—no extra ETL code.

6 Table inventory (9 physical, ≤ 5 GB with Pagila sample)

Layer	Physical table	Rows (Pagila)	Comment
Bronze	br_customer_cdc	~5K	append only
	br_rental_cdc	~16K	
	br_payment_cdc	~15K	
Silver	dim_customer_hist	~599	system-versioned
	dim_film_hist	~1K	system versioned
	fact_rental	~16K	partition by month
	fact_payment	~15K	partition by month
Gold	(materialised) vw_monthly_revenue	84	7 stores × 12 months
	(view) vw_top_titles	1K	

Everything else (dates, store dimension, etc.) can be small lookup views.

(The raw Pagila DDL is in your repo for reference.) [pagila-schema](#)

7 Putting it in the repo

pgsql [Copy](#)

```
pagila-medallion/
├── 00_ddl/           original OLTP DDL
├── 10_bronze_ingest/
│   ├── create_bronze.sql
│   └── load_from_wal.py
├── 20_silver_models/
│   ├── dim_customer_hist.sql
│   ├── fact_rental.sql
│   └── .
├── 30_gold_marts/
│   ├── vw_monthly_revenue.sql
│   └── .
└── dags/             Airflow / Dagster job definitions
└── docs/
    └── architecture.md
```

Tips

- Write a dbt macro `scd_system_versioned(table_name)` so every new dimension gets temporal logic for free.
- Add a notebook titled “Time-travel with SQL”—it’s a 5-minute wow moment for stakeholders.

8 Key take-aways for reviewers

1. Everything is reproducible – bronze raw files + deterministic SQL give identical results on every clone.
2. Time-travel comes “for free” – because the temporal pattern sits in silver, not bolted onto gold dashboards.
3. Small but realistic – Pagila’s relational quirks (partitioned payments, FK cascades, text search, etc.) illustrate challenges people face in real warehouses, yet stay <10MB.